

Homework 8.

Due on Gradescope by 5 PM, Saturday, June 5

Submission format. Please submit Google Colab your notebook in **.ipynb** format on Gradescope. Students are recommended to organize their code into modular cells with proper comments for ease of readability.

1. **Training RNNs with different regularization techniques:** Consider the simple RNN tutorial provided by PyTorch: https://pytorch.org/tutorials/intermediate/char_rnn_generation_tutorial.html. Follow this tutorial to first code up a simple RNN for generating names, given a language and starting alphabet as input. The estimated time to train a single model is 6 minutes, and you may need to train 5 or more models. You will need to begin by entering the following commands in your colab notebook to download the data:

```
!wget https://download.pytorch.org/tutorial/data.zip
!unzip data.zip -d
```

The training data that you download from the above link would contain names in 18 languages. Before starting the training:

- Split the provided data into train and validation set: choose the last 15% of names from each category as validation set (e.g., you can do it by creating a new dictionary *category_lines_val* similarly as *category_lines*, but with last 15% of entries from *lines* variable. You will be using this set to observe differences across various experiments by plotting training curves (loss vs. iterations). Note: you will need to randomly shuffle the *lines* variable in “Preparing the Data” section before assigning a subset of *lines* to validation set.
- Add a function to evaluate the performance of trained models on the entire validation set and output the average negative log-likelihood (this is the loss being used in the PyTorch tutorial linked above).
- Modify the main training loop to run your validation function every 5000 iterations.
- At the very beginning of your code, set the random seed to 4 manually using `torch.manual_seed(4)`, `numpy.random.seed(4)`, and `random.seed(4)`.

Your tasks for this problem are to train the given RNN with different regularization techniques and report the training behavior observed. Specifically:

- (a) The original PyTorch code uses dropout. Begin by training your network with this default setting and show the plot for average validation loss vs. number of iterations.

- (b) Next, repeat step (a) but *without* dropout.
- (c) Repeat step (a) but this time with L_2 regularization (also known as weight decay) instead of dropout. Note that you will need to search for the optimal weight for the L_2 regularization term. Do the plot as in step (a) for this experiment, after finding the most optimal weight for L_2 regularization.
- (d) Lastly, for all three experiments, generate names for the trained models by providing first letter of **Your First Name** for the following language categories: English, Portuguese, Russian, Greek, Chinese, Dutch.